

The k-ary LR method

It turns out that the LR method of repeated squaring can be generalized. Instead of breaking the exponent into bits of its base-2 representation, we can break it into larger pieces, and save some computations this way.

I'll present the k-ary LR method that breaks the exponent into its "digits" in base $m = 2^k$ for some integer k . The exponent can be written as:

$$b = t_i m^i + t_{i-1} m^{i-1} + \dots + t_0 m^0$$

Where t_i are the digits of b in base m . a^b is then:

$$a^{t_i m^i} \cdot a^{t_{i-1} m^{i-1}} \cdot \dots \cdot a^{t_0}$$

We compute this iteratively as follows [6]:

Raise a^{t_0} to the m -th power and multiply by a^{t_1} . We get $r_1 = a^{t_0 m + t_1}$. Next, raise r_1 to the m -th power and multiply by a^{t_2} , obtaining $r_2 = a^{t_0 m^2 + t_1 m + t_2}$. If we continue with this, we'll eventually get a^b .

This translates into the following code:

```
def modexp_lr_k_ary(a, b, n, k=5):
    """ Compute a ** b (mod n)

    K-ary LR method, with a customizable 'k'.
    """
    base = 2 << (k - 1)

    # Precompute the table of exponents
    table = [1] * base
    for i in xrange(1, base):
        table[i] = table[i - 1] * a % n

    # Just like the binary LR method, just with a
    # different base
    #
    r = 1
    for digit in reversed(_digits_of_n(b, base)):
        for i in xrange(k):
            r = r * r % n
        if digit:
            r = r * table[digit] % n

    return r
```

Note that we save some time by pre-computing the powers of a for exponents that can be digits in base m . Also, the `_digits_of_n` is the following generalization of `_bits_of_n`:

```
def _digits_of_n(n, b):  
    """ Return the list of the digits in the base 'b'  
        representation of n, from LSB to MSB  
    """  
    digits = []  
  
    while n:  
        digits.append(n % b)  
        n //= b  
  
    return digits
```

Performance of the k-ary method

In my tests, the k-ary LR method with $k = 5$ is about 25% faster than the binary LR method, and is within 20% of the built-in `pow` function.

Experimenting with the value of k affects these results, but 5 seems to be a good value that produce the best performance in most cases. This is probably why it's also used as the value of k in the implementation of `pow`.